




Training and Improving Neural Networks

How to train your neural network...
so that it doesn't explode

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com





sli.do

#DeepLearning

Table of Contents

- Regularization
- Bias and variance
 - Error analysis
- Optimization algorithms
- Hyperparameter tuning
- Normalization

Bias and Variance

Machine learning practices
using big(ger) data

Regularization

- Usual L1 and L2 rules apply


```
from tensorflow.layers import Dense
from tensorflow.keras import regularizers

Dense(
    kernel_regularizer = regularizers.L1L2(l1 = 0.5, l2 = 1),
    bias_regularizer = regularizers.L1L2(l1 = 0.5, l2 = 1),
    activity_regularizer = regularizers.L1L2(l1 = 0.3, l2 = 10))
```

- Regularization is applied to the **loss function**
 - It tries to "remove" or shrink the parameters
- We can regularize weights, biases and outputs
 - Usual steps: same regularization for weights and biases, none for outputs
- Note: using ReLU may result in activations = 0
 - This produces "dead neurons"
 - May be used as a form of regularization

Dropout

- Select a layer l
- At each training step, set a random fraction p of input weights of layer l to 0 \Rightarrow keep $1 - p$ units
 - To keep the dimensions, scale the remaining weights by $\frac{1}{1-p}$



```
from tensorflow.layers import Dropout  
Dropout(0.1)
```

- **Don't apply dropout during inference!**
 - tensorflow takes care of this

Selecting and Splitting Data

- Usually, we split the dataset like this
 - Training set – 70%
 - "Real training" set – 63%, validation set – 7%; 10 times
 - Testing set – 30%
- With many samples, this is unnecessary
 - And time consuming
- Law of big numbers
 - We can get stable results with many samples
 - \Rightarrow we have less chance of variance due to a small sample size
- Usual splitting for big data (e.g., 1M samples)
 - 980 000 / 10 000 / 10 000 samples
 - Alternatively, a bigger validation set: 980 000 / 16 000 / 4 000

Bias-Variance Error Analysis

- Bayes optimal error: the "real" error in data
 - No way to calculate, we need to try to come up with a measure
 - Naïve: this is 0%, the dataset is perfect
- Example: two-class classification (cats vs. dogs)
 - Metric: misclassification error ($E = 1 - A$)
 - Humans can achieve 0,5% error





Algorithm	Train set error	Validation set error	Bias, %	Variance, %	Verdict
A1	1%	11%	0,5%	10%	High variance
A2	15%	16%	14,5%	1%	High bias
A3	15%	30%	14,5%	15%	Both
A4	0,5%	1%	0,5%	0,5%	"Neither"
A5	0,3%	0,4%	?	?	?

Taking the Next Step

- There are no set rules, only things we can try
- High bias
 - Train a bigger network
 - Possibly, try out different architectures
 - Try to find one which is best suited for the task
 - Train longer (e.g., more **epochs**)
- High variance
 - Apply regularization
 - Try a smaller network architecture
 - Get more data
 - Or try to augment the current dataset
 - E.g. bootstrap sampling, image rotation, adding noise, etc.

Training / Validation Curves

- The same as what we already know
 - Plot a metric (e.g., loss, accuracy...) w.r.t. the dataset size or epoch
- The shape and relative position of both curves help diagnose under- / overfitting



Optimization

"Learn smarter, not harder"

Weight Initialization

- Vanishing / Exploding gradients problem
 - Deeper networks can learn very complex functions
 - \Rightarrow more layers = better
 - But let's look at what a computation looks like
 - Take, for example the activation at the 15th layer
 - Ignoring the activation functions for simplicity
 - $a^{[15]} \approx w^{[14]}a^{[14]} \approx w^{[14]}w^{[13]}a^{[13]} \approx \dots \approx w^{[0]}w^{[1]} \dots w^{[14]}x$
 - If the weights are similarly scaled, the product becomes $\approx w^{15}$
 - If some elements of w are $\gtrsim 1$, the product will become **really big**
 - Alternatively, if some elements are $\lesssim 1$, the product will become **really small**
 - This leads to problems when updating weights: $w = w - \nabla w$
 - The gradients either become $\approx \infty$, or ≈ 0
- Solution: **initialize the weights properly**

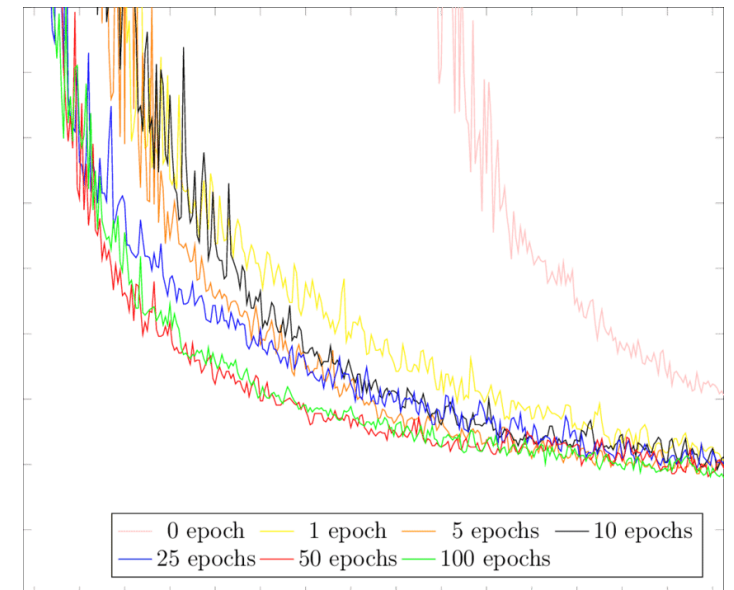
Weight Initialization (2)

- First, we know that we need random initialization
 - Gaussian, $\mu = 0, \sigma = 1$
 - $\mu = 0$ is needed because any bias has already been accounted for
- Also, initialize the weights with small numbers
 - The exploding / vanishing gradient problem affects only the first stages of training
 - After that, the NN should learn proper weights
- Glorot (Xavier) initialization
 - $\text{init} \sim N(0, \sigma)$ where $\sigma = \sqrt{\frac{2}{n_{in} + n_{out}}}$
 - where n_{in} and n_{out} are the numbers of input and output units of the layer

```
Dense(  
    kernel_initializer = tf.glorot_normal_initializer(), # or None  
    bias_initializer = tf.zeros_initializer())
```

Mini-batch Gradient Descent

- It takes a lot of time to pass through the entire dataset to perform only 1 step of GD (**batch gradient descent**)
 - Solution: take a random sample (**mini-batch**) each time: **mini-batch gradient descent**
 - If the mini-batch contains one sample \Rightarrow **stochastic GD** (SGD)
- The cost function will not decrease smoothly
 - But will tend to decrease, also the training will be faster
- Choosing a mini-batch size (n_b)
 - Powers of 2 lead to better speed
 - E.g., **32**, **64**, 128
 - Implementation
 - Shuffle the training set
 - At each training step, pass n_b examples



Improving Gradient Descent

- Momentum

- When updating weights, a fraction β_1 of the previous vector is added to the current:

$$v_t = \beta_1 v_{t-1} + (1 - \beta_1) \nabla J$$

$$w_t = w_{t-1} - \alpha v_t$$

- This tends to average out the steps in the "wrong" direction and speed up convergence



- RMSprop

- Similar to momentum, but second-order

$$S_t = \beta_2 S_{t-1} + (1 - \beta_2) (\nabla J)^2$$

$$w_t = w_{t-1} - \alpha \nabla J / (\sqrt{S_t} + \varepsilon)$$



Adam Optimizer

- Adam (**A**daptive **M**oment Estimation)
 - Combines momentum and RMSprop
 - Usually: α (tuning); $\beta_1 = 0,9$; $\beta_2 = 0,999$; $\varepsilon = 10^{-6} - 10^{-8}$
- Usage
 - In place of GradientDescentOptimizer
 - It's best to tune all hyperparameters but the default ones should work for most cases
 - Tuning α is non-negotiable

```
tf.train.AdamOptimizer(  
    learning_rate = 0.001,  
    beta1 = 0.9,  
    beta2 = 0.999,  
    epsilon = 1e-8)
```


A Note on Local Minima

- When training a model, GD and similar algorithms may get stuck in a local minimum
 - ML solution: different starting points
- In higher-dimensional spaces, most points with zero gradient are not local minima
 - They're instead saddle points
 - "Min" at one direction, "max" at the other
 - Example: 100 dimensions
 - Local min: all dimensions must be min
 - E.g., $p(\text{local min}) \approx 2^{-100} \approx 7,89 \cdot 10^{-31}$
 - When an optimizer gets to a saddle point, it's able to "roll off"






Hyperparameter Tuning

Similar to "standard"
machine learning

Prioritizing Hyperparameters

- Most important: learning rate α
- Momentum term β_1 , mini-batch size n_b
- Number of hidden units
- Number of hidden layers
- Search methodology
 - Grid search doesn't work (too large search space)
 - Use random search instead
- Better idea: use a coarse random search first
 - When you find a good place, "zoom in" to that
 - Repeat



Hyperparameter Scales

- Usual: uniform scale
 - E.g., hidden layers = {2, 3, 4}, hidden units $\in [50; 100]$
- Logarithmic scale
 - E.g., $\alpha \in [0,00001; 10]$
 - If we pick uniformly, most values will be close to 1
 - Solution: use a log scale for better search space exploration
 - $\alpha = 10^k, k \in [-5; 1]$
- Exponentially weighted averages (β_1, β_2)
 - E.g., $\beta \in [0,9; 0,9999]$
 - ⇒ $1 - \beta \in [0,1; 0,0001]$
 - ⇒ $1 - \beta = 10^k, k \in [-4; -1]$
 - $\beta = 1 - 10^k, k \in [-4; -1]$

Batch Normalization

- Normalizing inputs: Z-score

- $x = \frac{x - \mu}{\sigma}, \mu = \frac{1}{n} \sum_{i=1}^n x_i, \sigma^2 = \frac{1}{n-1} \sum_{i=1}^n x_i^2$

- Batch normalization

- At a given layer l , $z_n = \frac{z - \mu}{\sqrt{\sigma^2 + \epsilon}}$
- Use a linear transformation $\tilde{z} = \gamma z_n + \beta$ instead of the z
 - γ and β are parameters
 - γ and β are updated along with the weights w
- Application: compute **before** activation function
- Why does it work?
 - Doesn't allow the values to vary too much
- Implementation

```
from tensorflow.keras.layers import BatchNormalization
BatchNormalization(input)
```

Summary

- Regularization
- Bias and variance
 - Error analysis
- Optimization algorithms
- Hyperparameter tuning
- Normalization

The image features a white background with two blue decorative bars. The top bar is a solid blue strip. The bottom bar is a gradient blue strip that transitions from a lighter blue on the left to a darker blue on the right. The word "Questions?" is centered in a blue, sans-serif font.

Questions?