

Neural Networks for Language Processing

Thinking about language (and time series)

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com





sli.do

#DeepLearning

Table of Contents

- Time-dependent (sequential) models
 - Architecture
 - Types
- Improvements
- Word (token) representations
- Refinement algorithms
 - Attention
 - Transformers


Time-Dependent Models

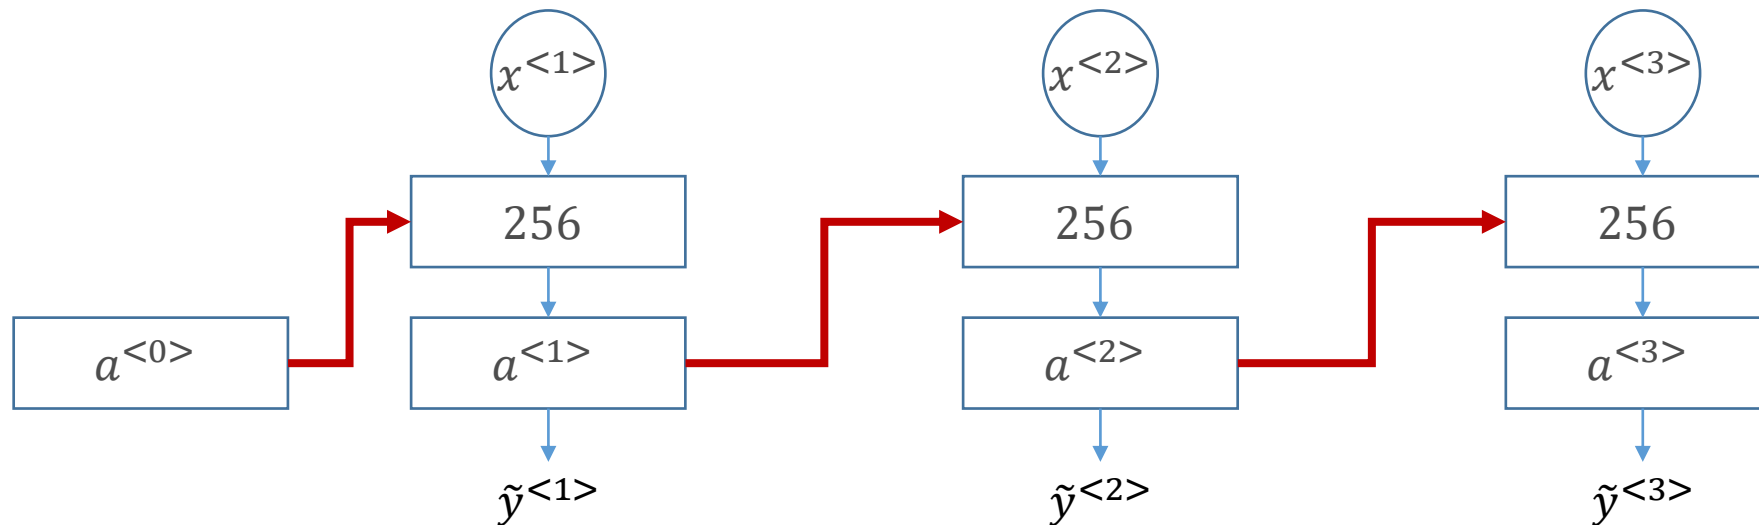
"Time is precious, waste it wisely"

Time-Dependent Model Examples

- Speech recognition: audio → transcript
- Machine translation: text (EN) → text (FR)
- Activity recognition: video → activity type (e.g. walking)
- Sentiment analysis: text → sentiment
- Generation
 - Text summarization
 - Music generation
- More generally, models whose inputs depend on time
 - "Standard" models: $\tilde{y} = f(x)$; recurrent models: $\tilde{y} = f(x, s)$
 - s – current state
 - Standard models don't allow variable-length inputs
 - Most standard models don't allow for weight sharing


Working with Sequences

- Training example: $x = \text{"A black cat in a box"}$
 - Split words (**tokenize** the input)
 - Present words as 1-hot encoded vectors using a dictionary (**vocabulary**)
 - $x^{<1>} = \text{"a"} = [1 \ 0 \ \dots \ 0]^T \equiv V_1$
 - $x^{<2>} = \text{"black"} = [0 \ 0 \ \dots \ 1 \ \dots \ 0]^T \equiv V_{329}$, etc.
 - Take a standard model (1-layer NN), pass each word
- 
- A black cat with yellow eyes is sitting inside a cardboard box, looking out. The box is on a light-colored carpet.



Recurrent Neural Networks



- "RNN cell"



- Deep architectures are also possible


- Learning: backpropagation through time

- The same as in a multi-layer network: $\arg \min J(y^{<t>}, \tilde{y}^{<t>})$



RNN Architectures

- One to one: standard
- One to many
 - Sequence generation given seed (e.g., image captioning)
- Many to one
 - One output for sequence (e.g., sentiment analysis)
- Many to many
 - Encoding and decoding (e.g., machine translation)
 - Synchronized output (e.g., video classification for each frame)



Language Model

■ Training

- Tokenize the input $x = [x^{<1>}; x^{<2>}; x^{<T_x>}]$
- Use a standard RNN, with no initial seed
 - $a^{<0>} = [0 \ 0 \ \dots 0] = \vec{0}, x^{<0>} = \vec{0}$
 - Output: \tilde{y} : a vector of probabilities for each word [0,0385 0,0476 ... 0,00041]
 - Softmax, with 10 000 outputs

■ Explanation

- First token: $\tilde{y} = P(w_1)$
- Second token: $\tilde{y} = P(w_2|w_1)$
- In general: $\tilde{y} = P(w_k|w_1, w_2, \dots, w_{k-1})$

■ Generation: random sampling according to computed P

- Input $x^{<0>} = \vec{0}, a^{<0>} = \vec{0}$; compute $a^{<1>}, \tilde{y}^{<1>}$; choose a word w_1
- Input $x^{<1>} \equiv w_1, a^{<1>}$; compute $a^{<2>}, \tilde{y}^{<2>}$; choose a word w_2
- ... until you reach [.]




Improved Models

Making things even more difficult

Vanishing Gradients

- RNN with a long input is similar to a very deep NN
- Examples
 - The match was long, but we won it which made us happy.
 - We **decided to go to the movies**, but our friend, who doesn't like scary movies, **didn't want to go**.
- Solution: Gated recurrent unit (GRU) – [Choi et al., 2014](#)
 - Update gate (z_t): how much of the past information to retain
 - Reset gate (r_t): how much information to forget
 - Final memory: current information + previous "context"



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$


$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Long-Short Term Memory (LSTM)


- Even more powerful (and complicated)
 - [Hochreiter and Schmidhuber, 1997](#)
 - This is only one layer, LSTM layers can also be stacked



- Basic parts of the architecture
 - Forget gate f_t
 - Update gate i_t
 - Cell state C_t
 - Output o_t
- A good [article](#) explaining LSTM cells

Bi-Directional Networks

- Intuition
 - RNNs may need information ahead, "from the future"
 - E.g. to translate word $w^{<5>}$ we may need the whole sentence
- Solution: just create pairs of networks



- These can be RNN, GRU, LSTM or other layers
- To compute activations, go left to right, then right to left

Representing Words

Find your way
in the multi-dimensional space

Word Representation

- Basic idea: one-hot encoding
- How to get insights on word relations?
 - Try to estimate **word features**: vectors of numbers for each word
 - Unsupervised process
 - **Embedding** from a space with one dimension per word to a lower-dimensional space (e.g. 300D)
 - Example uses
 - Use similarity measures (e.g. cosine distance) between vectors
 - Use projections to generate analogies ([Mikolov et al., 2013](#))
- Visualization: usually **t-SNE** or PCA
- Tensorboard uses Google's Embedding Projector
 - <https://www.tensorflow.org/guide/embedding>

Word2Vec and GloVe

- What we already described
 - A matrix E where each vocabulary word has a dense vector
- Context-target word pairs
 - Compute vectors for context and target
 - Loss: cross-entropy
- Similarity
 - Cosine similarity; closest words to "Sweden"
- Associations
 - Rome : Italy :: Beijing : China
 - king : queen :: man : woman
 - [Other examples](#)

| Word | Cosine distance |
|-------------|-----------------|
| norway | 0.760124 |
| denmark | 0.715460 |
| finland | 0.620022 |
| switzerland | 0.588132 |
| belgium | 0.585835 |
| netherlands | 0.574631 |
| iceland | 0.562368 |
| estonia | 0.547621 |
| slovenia | 0.531408 |

Refinement Algorithms

Some more tricks up our sleeves

Beam Search

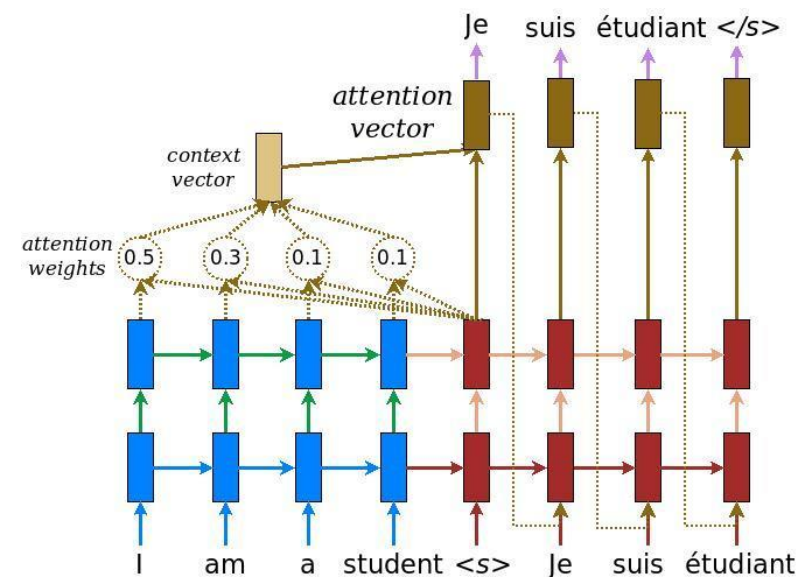

- Translation
 - Similar to generation, $\tilde{y} = f(x)$, maximize $P(\tilde{y}|x)$
- What if we have multiple candidates?
 - Use the language model to compute P
- Slight complications
 - *I am visiting NY this year end.*
 - *I am going to be visiting NY this year end.*
 - $P(\text{going}|\text{i, am}) > P(\text{visiting}|\text{i, am})$
 - Observations
 - One word at a time doesn't work too well
 - All words will require enormous computation power
 - Solution: **Beam search**
 - At each step, choose top B words (**beam width**)
 - More [details](#)

Attention

- [Xu et al., 2015](#)
- Another mechanism for dealing with complicated inputs
 - Another caveat: longer sentences have inherently lower probabilities so models tend to favor short sentences
 - Intuition: we don't need to know the entire sequence in order to be able to translate
- Idea
 - Use a bi-directional RNN (or GRU / LSTM)
 - For each part of the input $x^{<t>}$, compute "how much you care" about it: *attention*^{<t>}
- [Usages](#)
 - Translation, image captioning, speech recognition, text summarization, etc.

Attention Explained

- Queries, keys, values
 - Linear (dense) layers
- Multi-head attention
 - Just do the same operation multiple times 😊
- Sparse attention
 - Don't keep the whole attention table
- Scaled dot product attention
 - $A(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$
 - d_k - dimension of q and k
 - Outputs have variance = 1



Transformers

- [Brown et al., 2020](#) (GPT-3)
- Main points
 - Positional encoding
 - Attention blocks (heads) / self-attention
 - Encoder / decoder structure
- Usages
 - Language models, question answering, classification, paraphrasing / summarization, etc.
- Open issues
 - Attention blocks require too much memory
 - Too long training time

Summary

- Time-dependent (sequential) models
 - Architecture
 - Types
- Improvements
- Word (token) representations
- Refinement algorithms
 - Attention
 - Transformers

The image features a white background with two thick, wavy blue bars at the top and bottom. The top bar is a lighter blue, while the bottom bar is a darker blue. Centered on the white background is the word "Questions?" in a large, blue, sans-serif font.

Questions?