


Generative Models

Generating novel data, adversarial training

Yordan Darakchiev

Technical Trainer

iordan93@gmail.com





sli.do

#DeepLearning

Table of Contents

- Generative models: basics
- Unsupervised data generation
 - Style transfer
 - Generating sequences
 - Variational autoencoders
- Advanced generative models
 - Reminder: transformers
 - Generative adversarial networks
 - Stable diffusion
 - NeRF



Generative Models

Laying out the basics

Generative Models Overview

- Supervised learning: $(x, y) \Rightarrow \tilde{y} = f(x): y \approx \tilde{y}$
- Unsupervised learning: $x \Rightarrow f(x): f$ describes structure
- Generative models intuition
 - Learn the distribution of training data $p_{train}(x) = f(x)$
 - Make the model output new data $p_{model}(x)$ "similar" to p_{train}
 - In other words: after training, make a model: $\tilde{y} = f(\vec{0})$
- Examples
 - [PixelRNN / PixelCNN](#) – image completion
 - [DCGAN Image colorization](#)
 - [SRGAN](#) – "super resolution"
 - Sequence generation: words, [music](#), etc.; seq2seq models
 - Time series \Rightarrow planning (reinforcement learning)



Style Transfer

- [Gatys et al., 2016](#)
- [A tutorial](#)
- Objective
 - Input
 - Content image C , style image S
 - Output: image G with the content of C and style of S
- Start with a trained model (e.g. VGG-19) and train it further
- Loss function
 - Content loss L_c
 - Style loss L_s
 - Total loss: $L_{total} = \alpha L_c + \beta L_s$
 - α, β – numbers (relative contributions of the content and style)




Style Transfer (2)

- Content loss function
 - Extract activations from some middle layer l
 - Intuition
 - First layer: will force G to have very similar pixels to C
 - Last layer: "if there's a dog in C , draw a dog somewhere in G "
 - Therefore, it's better to choose a layer in between
 - Even better, treat l as a hyperparameter and try out a few architectures
- $L_{content}(C, G, l) = MSE(C^{[l]}, G^{[l]})$
 - C – original image, G – generated image
 - Pass both C and G to the trained network
 - You may not propagate them fully to save time
 - Take the output activations from layer l : $C^{[l]}, G^{[l]}$
 - These are 3D volumes
 - Apply MSE

Style Transfer (3)

- Style loss function
- How do we define style?
 - Correlations between different filter responses

- If we look at the activations at layer l , they will have the same style if they are correlated
 - Style matrix (Gram matrix): all dot products



- Diagonal G_{ii} : how powerful each filter is
 - Other $G_{ij}, i \neq j$: how correlated filters i and j are
 - Compute style matrices for G and S , minimize the distance
 - $L_s(S, G, l) = \text{MSE}(\text{Style}_G^{[l]}, \text{Style}_S^{[l]})$
 - Combine for many layers (with weights λ)
 - $L_{\text{style}}(S, G) = \sum_l \lambda_l \text{MSE}(\text{Style}_G^{[l]}, \text{Style}_S^{[l]})$; $\sum_l \lambda_l = 1$

Style Transfer (4)

- Training
 - Load the content and style images (C, S)
 - Generate G randomly
 - Load the trained model (e.g. VGG-19)
 - Pass C, S, G through the model, compute total loss
- Minimize total loss: standard optimization problem
 - Freeze all layers in VGG-19
 - This will change **the pixel values of G** to minimize the loss, **NOT the weights** of the trained model
- Obtaining a result
 - Take the updated image G
- Style transfer in tensorflow
 - [An in-depth tutorial](#)

Sequence Generators

Sampling novel sequences

Sequence Generation


- We've already seen this
 - [Language modelling](#)
 - [Music generation, more music generation](#) (Magenta)
 - [Image completion \(PixelRNN\)](#)
- Training
 - As usual, train a one-to-many RNN model
 - Maximize the likelihood of input data (Bayes)
- Generation
 - Use a seed input (e.g. initial words, or a partial picture), predict next
 - Continue until the end (e.g. [.] / <EOS> character or until the picture is done)
- Drawback: generation can be slow
 - Especially if we augment it by using beam search or a similar algorithm

Variational Autoencoders

Autoencoders with a simple trick

Variational Autoencoders

- Train an AE; throw away E
- Try getting an output from D
 - Pass $l \equiv x = \vec{0}$ to the decoder
 - Won't generate different samples :(
 - Pass random values
 - Produces random noise at output :(
- How to create latent vectors so the AE works properly?
 - Constrain E to produce specific vectors, e.g. $l \sim N(0, 1)$
 - You have now created a **variational autoencoder** (VAE)
 - Throw away E
 - Sample $x \sim N(0, 1)$
 - D now "understands" the input even though it's completely random



Variational Autoencoders (2)


- Loss function
 - Generative loss: how well D reconstructs x from l : $\text{MSE}(x, \tilde{y})$
 - Constraint loss: how close l is to $N(0, 1)$: [KL divergence](#)
 - $L = \text{MSE}(x, \tilde{y}) + \text{KLD}(\tilde{y}, N(0, 1))$
- Results
 - [MNIST](#), one more [MNIST](#)
 - [Music](#) (Roberts et al., 2017)
 - [Time series \(RL\)](#) (Gregor and Besse, 2018)
- Problems regarding generation
 - [Distribution matching](#) (Rosca et al., 2018)
 - [Adding more constraints to VAEs](#) (Rezende and Viola, 2018)

Advanced Generative Models

Learning how to be the best

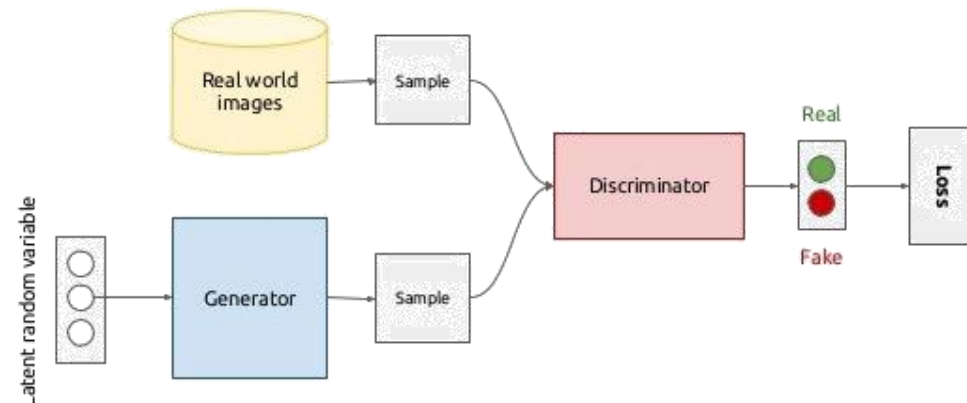
Adversarial Training

- [Goodfellow et al., 2014a](#)
- ML models misclassify examples that are only slightly different from correctly classified examples
 - "Overfitting"; many reasons
 - Look at the summary of the paper for a quick overview
 - Regularization alone doesn't help much
 - Adversarial training can serve as an additional regularizer
- Generating adversarial examples
 - Similar to style transfer (but much easier)
 - Pass a valid image and a noise G to the (frozen) neural network
 - Using the model gradients, update G to maximize the desired output




Generative Adversarial Networks

- [Goodfellow et al., 2014b](#)
 - One of the "cornerstone" papers of modern deep learning
- Approach: learn to generate from the training distribution by a two-player game
 - Instead of modelling vector densities (like we showed so far)
- High-level intuition
 - Generator network G : noise \Rightarrow output
 - Discriminator network D : input \Rightarrow is input fake?
 - Objective
 - Train together
 - D tries to distinguish between real and fake images
 - G tries to fool D



Generative Adversarial Networks (2)

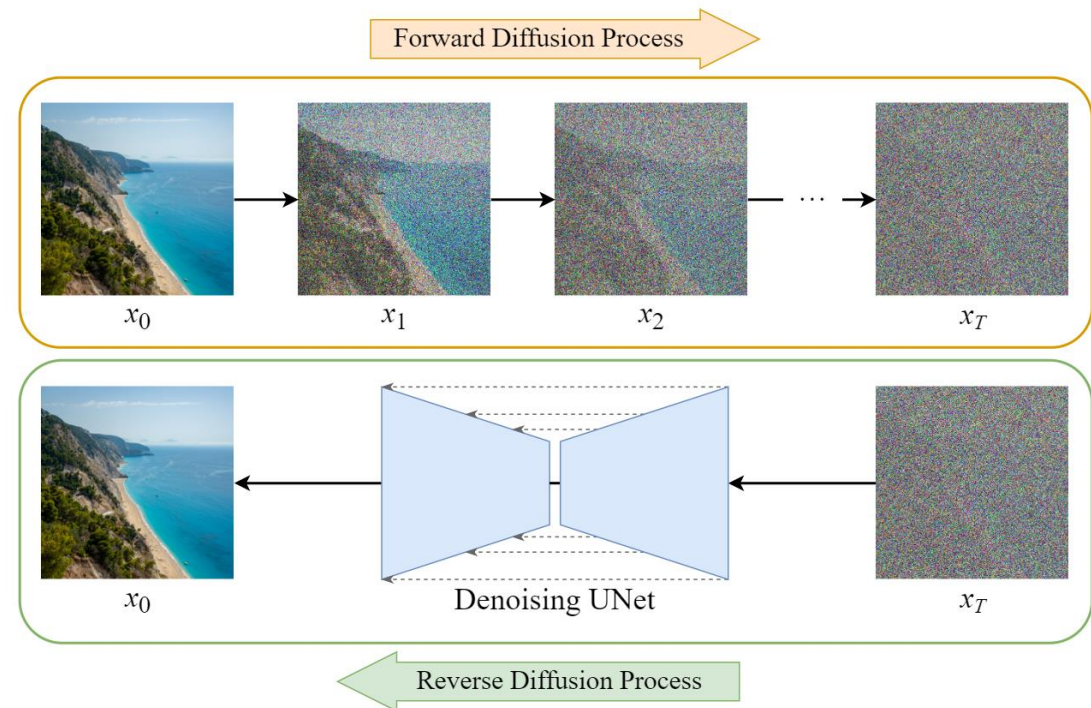
- After training, use G to generate outputs
 - Provide random input
 - Take generated output
- Another useful property
 - Learned features can be interpreted (similar to word2vec)



- GAN intuition
- Lots of examples
- Many more examples
- 17 hacks to make GANs work

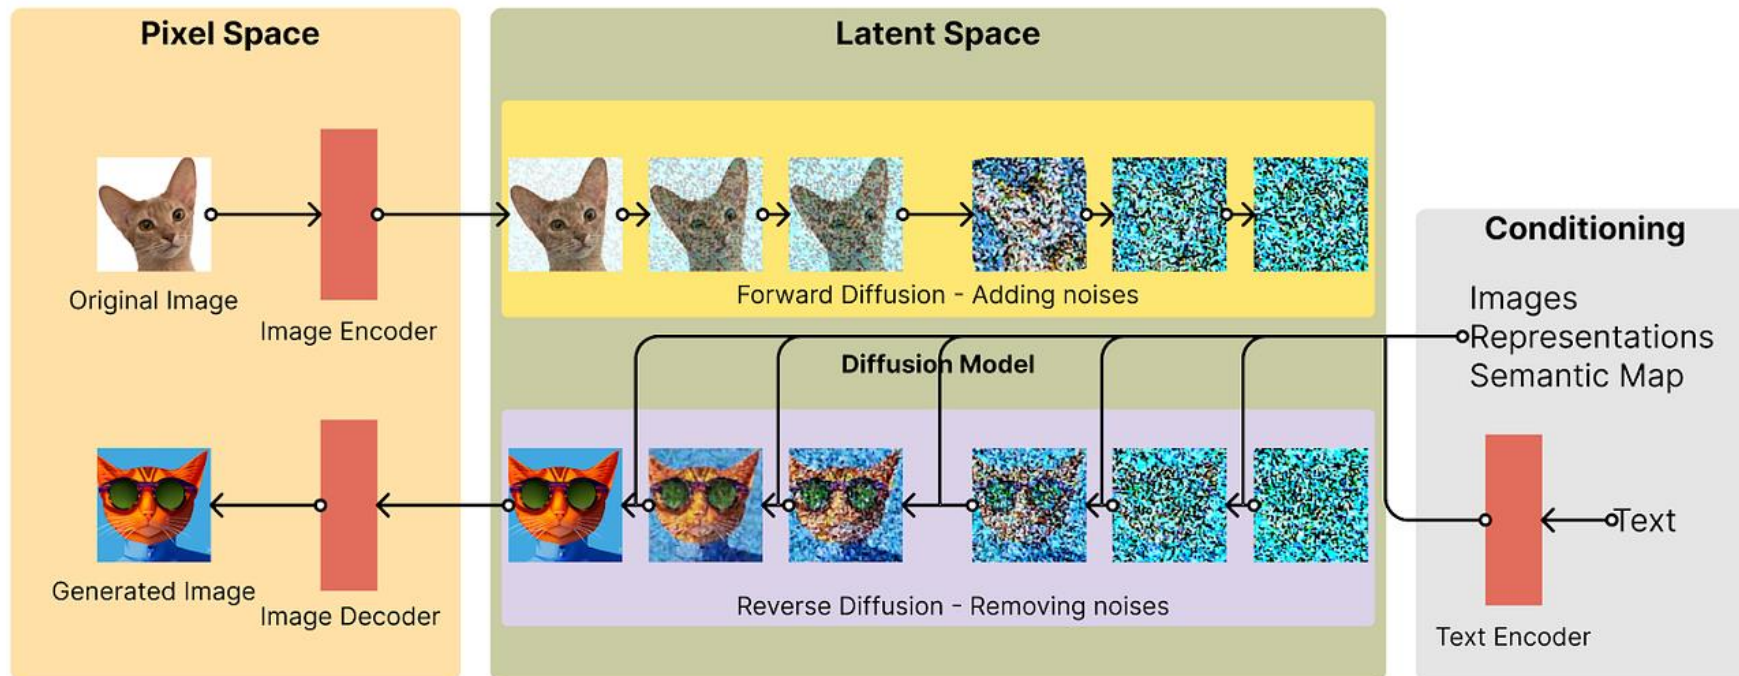
Stable Diffusion

- Inspiration: diffusion in physics
- Time-based process
 - Given noisy image $I(t) = I_{clear} + s(t) n$
 - Gaussian noise: $n \sim \mathcal{N}(\mu, \sigma)$
 - Predict $\tilde{n}(t)$ and subtract: $I(t - 1) = I(t) - \tilde{n}(t)$
- SD clearly explained



Stable Diffusion (2)

- Conditioning
 - Pass text prompt at the beginning
 - Prefer $\tilde{n}(t \mid \text{prompt})$ to $\tilde{n}(t \mid \text{no prompt})$



Summary

- Generative models: basics
- Unsupervised data generation
 - Style transfer
 - Generating sequences
 - Variational autoencoders
- Advanced generative models
 - Reminder: transformers
 - Generative adversarial networks
 - Stable diffusion
 - NeRF

The image features a white background with two decorative blue bars. The top bar is a solid blue strip. Below it is a thinner, darker blue wavy line. The bottom of the image is framed by a similar pattern of a thin dark blue wavy line above a solid blue strip.

Questions?